

HTTP/3 LAB



Bar Binyamin Varsulker,

Ori Meir Kushnir,

Preparation Questions

- a. What is the main difference between TCP and UDP?
- b. What are the advantages and disadvantages of using TCP versus UDP?
- c. What is QUIC? And on which protocol is it based on?
- d. What are the advantages of HTTP/3 compared to previous protocols (HTTP/1.1, HTTP/2)?
- e. What problems of HTTP/1.1 does HTTP/2 solve?
- f. What is Head-of-Line Blocking? How does it manifest in TCP?
- g. How does HTTP/3 address Head-of-Line Blocking?
- h. What is SCTP? What are the fields in the header?
- i. What is multi-homing and multi-streaming in SCTP?
- j. What are SCTP chunks?
- k. Give 2 advantages and 2 disadvantages of SCTP.
- l. What is a Docker container, and how is it different from a virtual machine?
- m. Why is Docker useful for running a simple HTTP server in a lab environment?

Part 1 - HTTP/1.1, HTTP/2 and HTTP/3 with Caddy Web Server

Step 1: Create an Ubuntu Virtual Machine

- You can use virtualization tools such as VirtualBox or VMware.
- Required OS: Ubuntu 22.04.
- Minimum Resources: 1 CPU, 2GB RAM, 5GB free disk space.
- Update the VM by running the following command in the terminal:
`sudo apt update && sudo apt upgrade -y`

Step 2: Install Docker on Ubuntu

- Install Docker engine:
`sudo apt install docker.io -y`
- Enable and start Docker service:
`sudo systemctl enable docker`
`sudo systemctl start docker`

Verify Docker is working:

```
docker --version
```

Step 3: Configure a Simple Caddy Server

- Create a working directory:
`mkdir ~/site`
`cd ~/site`
- Create the Caddyfile:
`nano Caddyfile`
- Paste the following content:

```
localhost:443 {  
    root * /srv
```

```
    file_server
    templates
    tls internal
}

localhost:80 {
    root * /srv
    file_server
    templates
}
```

Question 1: Explain each block in Caddyfile

- Create a webpage:

```
nano ~/site/index.html
```

- Paste the following configuration and save it:

```
<!DOCTYPE html>

<html>

<head><title>Connection Info</title></head>

<body>

    <h1>Hello lab students!</h1>

    <p><strong>Using protocol:</strong> {{.Req.Proto}}</p>

</body>

</html>
```

Step 4: Run the Caddy Server Using Docker

- Run the following command:

```
docker run -d \  
  -p 80:80 \  
  -p 443:443 \  
  -p 443:443/udp \  
  -v ~/site:/srv \  
  -v caddy_data:/data \  
  caddy:latest \  
  caddy run --config /srv/Caddyfile
```

Question 2: Explain the above command.

- Confirm that the container is running:
`sudo docker ps`

1. Send an HTTP/1.1 Request

- Install curl:
`sudo apt install curl -y`

Question 3: Explain about curl.

- Open Wireshark and start capture on lo (loopback) interface.
Use capture filters:
 - `tcp.port == 80` (HTTP/1.1 plaintext)
 - `tcp.port == 443` (HTTP/2)
 - `udp.port == 443` (HTTP/3 QUIC)
- Run the command:
`curl --http1.1 http://localhost`
- Stop the capture.

Question 4: Explain what is a major disadvantage in HTTP/1.1 that you can see in the Wireshark pcap and add supporting screenshots.

2. Send an HTTP/2 Request

- Start a new Wireshark capture and run the following command:
`curl --http1.1 http://localhost`

Question 5: Analyze HTTP/2 packet from the Wireshark pcap and add the packet to the final report. Compare It with the HTTP/1.1 packets from the previous question.

3. Send an HTTP/3 Request

- Start a new Wireshark capture and run the following command:
`curl --http3 -k --trace-ascii - https://localhost`
- You'll probably get an error which says that the currently installed curl version doesn't support HTTP/3. Follow instructions in the appendix and afterwards run this command again.

Question 6: Compare the output of the above command with the HTTP/3 packets from Wireshark. What differences can be observed?

Part 2 - HTTP/2: Demonstrating TCP-Level Head-of-Line Blocking

In this part of the lab, you will observe and understand the TCP-level Head-of-Line blocking issue inherent in HTTP/2.

HTTP/2 allows multiple streams to be multiplexed over a single TCP connection. However, TCP guarantees in-order delivery across all bytes in the connection. If a packet is lost during transmission, TCP must hold back all following packets — even those belonging to different HTTP/2 streams — until the lost packet is successfully retransmitted.

This behavior leads to head-of-line blocking: a delay in delivering data from one stream blocks the progress of other streams sharing the same TCP connection.

Later, we will see how HTTP/3 resolves this issue.

Step 1 - Preparing the Test Files

Before starting the HTTP/2 HOL blocking experiment, we need to create and set up the files that the server will serve: one **small file** and two **large files**.

- **Create the files on your VM:**

Run the following commands in ~/site directory to create the required files:

Create a small fast.html file (~100 bytes)

```
echo "<html><body><h1>Fast Response</h1></body></html>" > fast.html
```

Create a large file (~20 MB) for slow1.dat

```
dd if=/dev/urandom of=slow1.dat bs=1M count=20
```

Create a large file (~30 MB) for slow2.dat

```
dd if=/dev/urandom of=slow2.dat bs=1M count=30
```

Step 2 - Introduce Network Impairments

Introduce network impairments on the loopback interface using tc (traffic control):

```
sudo tc qdisc add dev lo root netem delay 100ms loss 15%
```

Step 3 – Download multiple files in parallel over HTTP/2 using curl

Run the following commands:

```
sudo curl --http2 -k --parallel --parallel-immediate \  
  --output slow1_downloaded.txt \  
  --output slow2_downloaded.txt \  
  --output fast_downloaded.txt \  
  --trace-ascii trace_http2.log \  
  --trace-time \  
  https://localhost/slow1.dat \  
  https://localhost/slow2.dat \  
  https://localhost/fast.html
```

```
https://localhost/fast.html
```

```
grep -E 'GET|HTTP/2' trace_http2.log
```

Question 7:

- i. **Explain these commands.**
- ii. **Analyze the output and explain why does the small file (fast.html) might arrive last, even though HTTP/2 allows multiplexed streams? How is this affected by TCP's design? Attach screenshots.**

Step 4 –HTTP/2 vs. HTTP/3 – Head-of-Line Blocking

In this step, you will experimentally observe the differences between HTTP/2 and HTTP/3 in terms of Head-of-Line blocking and understand how HTTP/3 solves a major limitation of TCP-based HTTP/2.

Follow the instructions carefully:

- Make sure your Caddy server is correctly set up to serve all three files from the previous step.
- Download the files over HTTP/3:

```
sudo curl --http3 -k --parallel --parallel-immediate \  
  --output slow1_downloaded.txt \  
  --output slow2_downloaded.txt \  
  --output fast_downloaded.txt \  
  --trace-ascii trace_http3.log \  
  --trace-time \  
  https://localhost/slow1.dat \  
  https://localhost/slow2.dat \  
  https://localhost/fast.html
```

Question 8: Compare the output log here to the output log from the previous step. Can you conclude that there is no TCP-Level HOL blocking? Add screenshots.

Part 4 – SCTP

Step 1 – iperf Server Setup

In this part you will run Docker containers for iperf server and client to test SCTP communication.

In a terminal window, run the following command to start a Docker container with iperf server:

```
sudo docker run --rm -t --net=host --name iperf-server  
sofianinho/iperf3:3.6-ubuntu18.04 -s --sctp
```

Step 2 – Test SCTP communication

1. Open Wireshark and start capture.
2. Apply filter: sctp.
3. In a new terminal window, run the following command:

```
sudo docker run --rm -t --name iperf-client sofianinho/iperf3:3.6-  
ubuntu18.04 -c <server-ip> --sctp
```

*server-ip can be obtained by running `hostname -I` (this is the host IP address).

4. Wait for the command to finish execution and then stop the capture.

Question 9: Compare the SCTP handshake to the TCP handshake that you are already familiar with and is used for instance in HTTP/2. What differences do you observe in terms of the number of steps? Which one includes cookie-based verification? Explain and add screenshots.

Question 10: Why do you not see HEARTBEAT packets in your capture?

Question 11: Compare SCTP to QUIC (used in HTTP/3). How does SCTP's multistreaming compare to QUIC's multiplexing?

Good Luck!

Appendix – Install curl with HTTP/3 Support

Follow these steps on Ubuntu to compile and install curl with HTTP/3 (QUIC) support alongside HTTP/1.1 and HTTP/2.

Step 1: Install Build Dependencies

```
sudo apt update
sudo apt install -y \
    git build-essential autoconf libtool pkg-config \
    libev-dev libssl-dev zlib1g-dev \
    cmake curl libbrotli-dev libnghttp2-dev \
    python3 python3-venv python3-pip
```

Step 2: Build and Install OpenSSL v3.5+

```
cd ~/Desktop
git clone --depth 1 -b openssl-3.5.0 https://github.com/openssl/openssl
cd openssl
./config --prefix=$HOME/opt/openssl --libdir=lib
make -j$(nproc)
make install
```

Step 3: Build and Install nghttp3

```
cd ~/Desktop
git clone -b v1.3.0 https://github.com/ngtcp2/nghttp3
cd nghttp3
git submodule update --init
autoreconf -fi
./configure --prefix=$HOME/opt/nghttp3 --enable-lib-only
make -j$(nproc)
make install
```

Step 4: Build and Install ngtcp2

```
cd ~/Desktop
git clone -b v1.3.0 https://github.com/ngtcp2/ngtcp2
cd ngtcp2
autoreconf -fi
./configure \
    PKG_CONFIG_PATH="$HOME/opt/openssl/lib/pkgconfig:$HOME/opt/nghttp3/ \
    lib/pkgconfig" \
    LDFLAGS="-Wl,-rpath,$HOME/opt/openssl/lib" \
    --prefix=$HOME/opt/ngtcp2 \
    --enable-lib-only \
    --with-openssl
make -j$(nproc)
make install
```

Step 5: Build and Install curl with HTTP/3 Support

```
cd ~/Desktop
git clone https://github.com/curl/curl
cd curl
autoreconf -fi
./configure \
    --with-ssl=$HOME/opt/openssl \
    --with-nghttp2 \
    --with-nghttp3=$HOME/opt/nghttp3 \
    --with-ngtcp2=$HOME/opt/ngtcp2 \
    --prefix=$HOME/opt/curl-http3 \
    LDFLAGS="-Wl,-rpath,$HOME/opt/openssl/lib"
make -j$(nproc)
make install
```

Step 6: Use the Compiled Version

Add it to your PATH:

```
export PATH=$HOME/opt/curl-http3/bin:$PATH
```

Verify:

```
curl --version
```

Expected output should include:

```
Features: ... HTTP1.1 HTTP2 HTTP3 ...
```