

HTTP/3 LAB



Lab Report

Bar Binyamin Varsulker,

Ori Meir Kushnir,

..

Answers to Practical Questions

1. The file defines two server blocks — one for HTTPS (localhost:443) and one for HTTP (localhost:80). Each block tells Caddy what to do when a request comes in on that port.

Starting with localhost:443, this is the configuration for handling secure HTTPS requests. Inside the block, we first set the root directory for the server to /srv, meaning all requests will look for files starting from that directory. Then, file_server tells Caddy to serve static files (like HTML, CSS, images, etc.) from that root folder. The templates directive allows simple templating, so if you're using placeholders like {{.Host}} in your HTML files, Caddy can process them. Finally, tls internal tells Caddy to automatically generate and use a TLS certificate from its internal CA, which is super helpful for local development — you get HTTPS without needing to get a real certificate.

The second block, localhost:80, is for regular HTTP traffic. It's basically doing the same as the first block — setting the root to /srv, enabling static file serving, and allowing template processing — but it doesn't include TLS because it's HTTP, not HTTPS.

This kind of setup is useful for running a basic secure web server locally, where HTTPS is supported on port 443 with automatic internal certificates, and HTTP works on port 80 for fallback or redirecting purposes.

2. This command runs a Caddy server inside a Docker container in detached mode, meaning it runs in the background.

The -p flags map ports from your host machine to the container. Port 80 is for HTTP, and port 443 is for HTTPS — it's mapped both for TCP (default) and explicitly for UDP as well. That UDP mapping is important because HTTP/3 (which Caddy supports) runs over QUIC, and QUIC uses UDP on port 443.

The -v flags mount volumes. The first one, ~/site:/srv, maps your local site folder into the container's /srv directory — that's where Caddy will find your website files and your Caddyfile configuration. The second volume, caddy_data:/data, is where Caddy stores things like TLS certificates and internal state, so they persist even if the container is restarted.

At the end, the command tells Docker to run the official Caddy image (caddy:latest) and start the server using caddy run, while explicitly pointing it to the configuration file located at /srv/Caddyfile, which is the one you mounted earlier.

3. curl (Client URL) is a CLI utility used for transferring data to or from servers using a variety of network protocols, including HTTP, HTTPS, FTP, SMTP, and others. It's commonly used by developers and system admins for downloading files, interacting with web APIs, debugging network connections, or automating web-related tasks. curl operates through simple commands entered directly into the terminal, allowing users to easily specify details like URLs, request methods, headers, and data payloads.

4.

The image shows a Wireshark packet capture of an HTTP 1.1 200 OK response. The packet list on the left shows packet 318 selected, which is an HTTP packet from 127.0.0.1 to 127.0.0.1. The packet details pane shows the 'Line-based text data: text/html (7 lines)' section expanded, displaying the raw HTML content. The packet bytes pane at the bottom shows the hex and ASCII representation of the data, which is the same HTML content as shown in the details pane.

No.	Time	Source	Destination	Protocol	Length	Info
316	252.830000662	127.0.0.1	127.0.0.1	HTTP	139	GET / HTTP/1.1
317	252.830004140	127.0.0.1	127.0.0.1	TCP	66	80 → 37022 [ACK] Seq=1
318	252.836255632	127.0.0.1	127.0.0.1	HTTP	381	HTTP/1.1 200 OK (text/html)
319	252.836259928	127.0.0.1	127.0.0.1	TCP	66	37022 → 80 [ACK] Seq=7
320	252.836400067	127.0.0.1	127.0.0.1	TCP	66	37022 → 80 [FIN, ACK]
321	252.836615184	127.0.0.1	127.0.0.1	TCP	66	80 → 37022 [FIN, ACK]
322	252.836619787	127.0.0.1	127.0.0.1	TCP	66	37022 → 80 [ACK] Seq=7

Line-based text data: text/html (7 lines)

```
<!DOCTYPE html>\n<head><title>Connection info</title></head>\n<body>\n<h1>Hello lab students!</h1>\n<p><strong>Using protocol:</strong>HTTP/1.1</p>\n</body>\n</html>\n
```

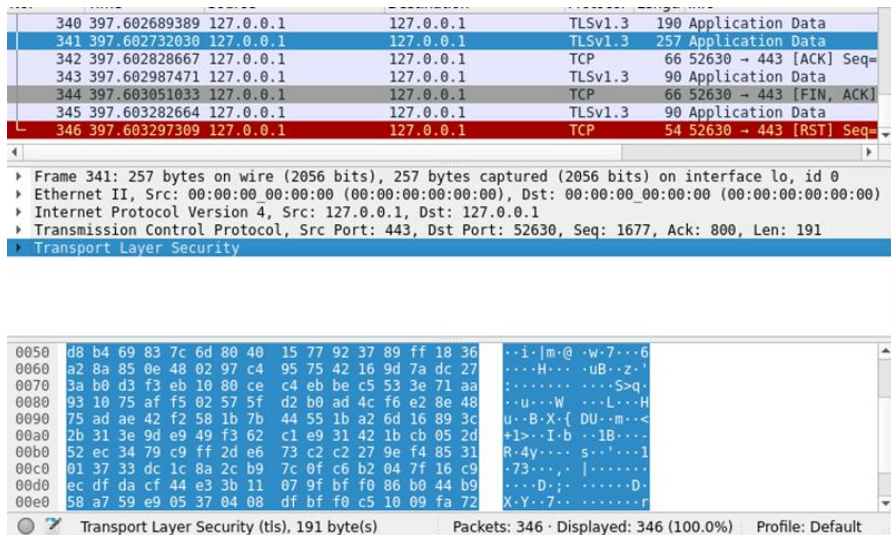
00f0 61 64 3e 3c 74 69 74 6c 65 3e 43 6f 6e 6e 65 63 ad<titl e>Connec
0100 74 69 6f 6e 20 69 6e 66 6f 3c 2f 74 69 74 6c 65 tion inf o</title
0110 3e 3c 2f 68 65 61 64 3e 0a 3c 62 6f 64 79 3e 0a ></head> <body>
0120 3c 68 31 3e 48 65 6c 6c 6f 20 6c 61 62 20 73 74 <h1>Hell o lab st
0130 75 64 65 6e 74 73 21 3c 2f 68 31 3e 0a 3c 70 3e udents!< /h1> <p>
0140 3c 73 74 72 6f 6e 67 3e 55 73 69 6e 67 20 70 72 Using pr
0150 6f 74 6f 63 6f 6c 3a 3c 2f 73 74 72 6f 6e 67 3e otocol:< /strong>
0160 48 54 54 50 2f 31 2e 31 3c 2f 70 3e 0a 3c 2f 62 HTTP/1.1 </p> </b
0170 6f 64 79 3e 0a 3c 2f 68 74 6d 6c 3e 0a ody> </h tml>

When analyzing the HTTP/1.1 traffic in Wireshark packet captures, the most glaring disadvantage becomes immediately apparent: all data is transmitted in plain text. This means anyone with access to network traffic can easily read usernames, passwords, session cookies, and other sensitive information directly from the pcap file without any special tools or decryption keys.

This Wireshark capture perfectly demonstrates the security vulnerability of HTTP/1.1 I mentioned. The image shows:

- A packet capture with HTTP/1.1 traffic (packet #318 shows "HTTP/1.1 200 OK")
- The HTML content of the web page is fully visible in plain text.
- Both the ASCII representation (right side) and hex values (left side) of the data are completely readable.

5.



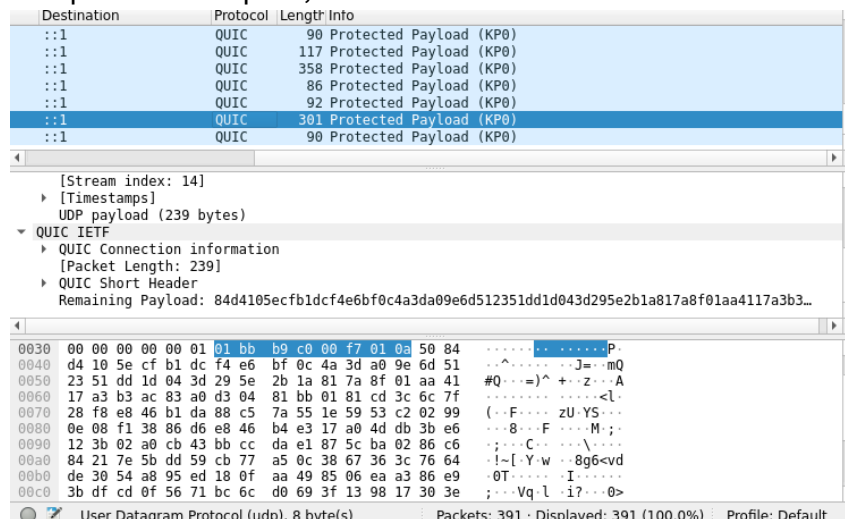
Looking at this Wireshark capture of HTTPS traffic using TLS, we can immediately see the huge difference compared to the previous HTTP/1.1 example.

This packet clearly shows encrypted data being transmitted - in the bottom section, all we can see is scrambled, unreadable bytes. The protocol column shows "TLSv1.3" instead of "HTTP," and the packet contents are labeled as "Application Data" rather than showing the actual HTTP content.

6.

```
<= Recv header, 23 bytes (0x17)
0000: vary: Accept-Encoding
<= Recv header, 37 bytes (0x25)
0000: date: Tue, 29 Apr 2025 16:38:34 GMT
<= Recv header, 40 bytes (0x28)
0000: content-type: text/html; charset=utf-8
<= Recv header, 2 bytes (0x2)
0000:
<= Recv data, 175 bytes (0xaf)
0000: <!DOCTYPE html>.<html>.<head><title>Connection Info</title></hea
0040: d>.<body>... <h1>Hello lab students!</h1>... <p><strong>Using pr
0080: otocol:</strong> HTTP/3.0</p>...</body>.</html>..
<!DOCTYPE html>
<html>
<head><title>Connection Info</title></head>
<body>
  <h1>Hello lab students!</h1>
  <p><strong>Using protocol:</strong> HTTP/3.0</p>
</body>
</html>
== Recv data, 0 bytes (0x0)
== Info: Connection #0 to host localhost left intact
ubu@master-node:~/site$
```

In the provided outputs, we observe clear



differences between using the curl command (with the **--trace-ascii** flag) and inspecting HTTP/3 traffic in Wireshark. The curl output directly shows readable details like HTTP/3 headers (such as content-type, content-length, and date) and the actual HTML content returned by the server. On the other hand, the Wireshark output only displays QUIC packets labeled as "Protected Payload," indicating that the data is

encrypted and not directly readable. This means curl gives an easy-to-understand view of the decrypted communication, whereas Wireshark captures the low-level, encrypted traffic, requiring additional steps or keys to interpret clearly.

In addition, curl explicitly mentions HTTP/3 protocol details, whereas Wireshark primarily identifies the underlying QUIC protocol.

7.

i. curl:

--http2: Force the client to use the HTTP/2 protocol when communicating with the server.

-k: Allow insecure SSL connections (skip certificate verification since we are using a local self-signed server).

--parallel --parallel-immediate: Fetch multiple URLs in parallel **immediately** instead of waiting for one to finish before starting the next.

--output slow1_downloaded.txt ...: Save each downloaded file to a specific output file name (one output per URL).

--trace-ascii trace_http2.log: Record a detailed ASCII trace of all sent/received HTTP/2 frames into trace_http2.log.

--trace-time: Add timestamps to every line inside the trace file, so the timing of events can be analyzed.

https://localhost/slow1.dat https://localhost/slow2.dat https://localhost/fast.html:
These are the three files being requested from the server over HTTPS on localhost.
This command downloads all three files at once over HTTP/2, records the entire session into a log file, and saves the files separately.

grep:

Searches inside the trace file for lines that contain either:

- the HTTP GET requests (showing what was requested),
- or the HTTP/2 responses (showing when a response was received).

-E: Use extended regular expressions so we can match multiple patterns (GET or HTTP/2).

This allows filtering only the important parts (requests and responses) from the trace.

ii.

```
ubu@master-node:~/Desktop/test$ grep -E 'GET|HTTP/2' trace_http2.log
23:00:15.219344 == Info: using HTTP/2
23:00:15.219429 == Info: [HTTP/2] [1] OPENED stream for https://localhost/slow1.dat
23:00:15.219434 == Info: [HTTP/2] [1] [:method: GET]
23:00:15.219437 == Info: [HTTP/2] [1] [:scheme: https]
23:00:15.219439 == Info: [HTTP/2] [1] [:authority: localhost]
23:00:15.219441 == Info: [HTTP/2] [1] [:path: /slow1.dat]
23:00:15.219443 == Info: [HTTP/2] [1] [user-agent: curl/8.14.0-DEV]
23:00:15.219445 == Info: [HTTP/2] [1] [accept: */*]
0000: GET /slow1.dat HTTP/2
23:00:15.420391 == Info: using HTTP/2
23:00:15.420425 == Info: [HTTP/2] [1] OPENED stream for https://localhost/slow2.dat
23:00:15.420428 == Info: [HTTP/2] [1] [:method: GET]
23:00:15.420430 == Info: [HTTP/2] [1] [:scheme: https]
23:00:15.420432 == Info: [HTTP/2] [1] [:authority: localhost]
23:00:15.420433 == Info: [HTTP/2] [1] [:path: /slow2.dat]
23:00:15.420435 == Info: [HTTP/2] [1] [user-agent: curl/8.14.0-DEV]
23:00:15.420437 == Info: [HTTP/2] [1] [accept: */*]
0000: GET /slow2.dat HTTP/2
0000: HTTP/2 200
0000: HTTP/2 200
23:00:15.886275 == Info: using HTTP/2
23:00:15.886303 == Info: [HTTP/2] [1] OPENED stream for https://localhost/fast.html
23:00:15.886305 == Info: [HTTP/2] [1] [:method: GET]
23:00:15.886307 == Info: [HTTP/2] [1] [:scheme: https]
23:00:15.886308 == Info: [HTTP/2] [1] [:authority: localhost]
23:00:15.886309 == Info: [HTTP/2] [1] [:path: /fast.html]
23:00:15.886310 == Info: [HTTP/2] [1] [user-agent: curl/8.14.0-DEV]
23:00:15.886312 == Info: [HTTP/2] [1] [accept: */*]
0000: GET /fast.html HTTP/2
0000: HTTP/2 200
0380: ...sM.C...1....Z....;''.GET-^.....T,%f.1=;.....&(Z...j.....
11c0: ~..GET...q..k..P..hKb...9.n..otqh..0...Z$.=-.....6....#..eb)...
ubu@master-node:~/Desktop/test$
```

Although HTTP/2 supports multiplexing multiple streams over a single TCP connection, it still relies on TCP for transport. TCP ensures strict in-order delivery of packets — meaning if one packet is lost or delayed, all following packets must wait until the missing packet is received.

In our experiment, the two large files (slow1.dat and slow2.dat) were requested first. Since they are big, they occupy a lot of TCP packets. Even though fast.html is very small and requested afterward, its packets are queued behind the packets of the large files at the TCP layer.

If even one packet from slow1.dat or slow2.dat is delayed or lost, TCP holds back all subsequent packets (including the ones for fast.html) until the missing packet is retransmitted and received. This causes fast.html to be delayed even though it could have been transmitted earlier if TCP allowed out-of-order delivery.

Thus, TCP-level HOL blocking prevents HTTP/2 from delivering small, fast responses quickly, despite HTTP/2's multiplexing features.

8.

```
23:56:42.965978 == Info: [HTTP/3] [0] [accept: /*/]
23:56:42.965989 => Send header, 84 bytes (0x54)
0000: GET /slow2.dat HTTP/3
0017: Host: localhost
0028: User-Agent: curl/8.14.0-DEV
0045: Accept: /*/*
0052:
23:56:42.966053 == Info: Request completely sent off
23:56:42.967738 <= Recv header, 13 bytes (0xd)
0000: HTTP/3 200
23:56:42.967770 <= Recv header, 15 bytes (0xf)
0000: server: Caddy
23:56:42.967776 <= Recv header, 23 bytes (0x17)
0000: vary: Accept-Encoding
23:56:42.967780 <= Recv header, 27 bytes (0x1b)
0000: etag: "d9iypefnq5fgiq8lc"
23:56:42.967785 <= Recv header, 46 bytes (0x2e)
0000: last-modified: Tue, 29 Apr 2025 07:59:23 GMT
23:56:42.967789 <= Recv header, 22 bytes (0x16)
0000: accept-ranges: bytes

735 0300: .j9...c06.y.4V0.....<...8U.F.N...r.*Y..8.U.l...tx.N.g...X
736 0340: .]J;....\..[!..K.m.^.....s...a.).../.g...w...v'.W~.-&
737 0380: L.s.p0hoc<...[,.....oE...8...:b.....U..x$.E.^..y.r....!u.Yd&.
738 03c0: ...`12..R...$86.kk1...L...R.b.....0.....I!..6.m...e.....
739 0400: .V;@...8.p.?s..9...[~%..j9_z'1.....V...J...bu...!....R.#..
740 0440: .py..E.?+!...1..h..R..3G...0p...0....Y..Jf.'.'...Y...i...
741 0480: .1.D.Q|y....9.e....|.8.....N.x$.:|].....n+..q..q...C...
742 04c0: sE.T.8.k...A.p.E...
743 23:56:42.971765 <= Recv header, 13 bytes (0xd)
744 0000: HTTP/3 200
745 23:56:42.972959 <= Recv header, 21 bytes (0x15)
746 0000: content-length: 124
747 23:56:42.972983 <= Recv header, 15 bytes (0xf)
748 0000: server: Caddy
749 23:56:42.972989 <= Recv header, 23 bytes (0x17)
750 0000: vary: Accept-Encoding
751 23:56:42.972994 <= Recv header, 37 bytes (0x25)
752 0000: date: Tue, 29 Apr 2025 20:56:42 GMT
753 23:56:42.972999 <= Recv header, 40 bytes (0x28)
754 0000: content-type: text/html; charset=utf-8

754 0000: content-type: text/html; charset=utf-8
755 23:56:42.973004 <= Recv header, 2 bytes (0x2)
756 0000:
757 23:56:42.973008 <= Recv data, 124 bytes (0x7c)
758 0000: <!DOCTYPE html>.<html>.<head><title>Fast</title></head>.<body>.<
759 0040: h1>Fast Response</h1>.<p>No delay here!</p>.</body>.</html>.<
760 23:56:42.973174 <= Recv data, 0 bytes (0x0)
761 23:56:42.973214 == Info: Connection #2 to host localhost left intact
762 23:56:42.976330 <= Recv header, 13 bytes (0xd)
763 0000: HTTP/3 200
764 23:56:42.976815 <= Recv header, 15 bytes (0xf)
765 0000: server: Caddy
766 23:56:42.976829 <= Recv header, 23 bytes (0x17)
767 0000: vary: Accept-Encoding
768 23:56:42.976834 <= Recv header, 27 bytes (0x1b)
769 0000: etag: "d9iypedjham4chhq8"
770 23:56:42.976838 <= Recv header, 46 bytes (0x2e)
771 0000: last-modified: Tue, 29 Apr 2025 07:59:23 GMT
772 23:56:42.976843 <= Recv header, 22 bytes (0x16)
773 0000: accept-ranges: bytes
```

In the provided trace logs, all three resources (slow1.dat, slow2.dat, fast.html) were requested in close succession and responded to with similar timestamps, regardless of their order. This demonstrates that HTTP/3 eliminates TCP-level Head-of-Line blocking, enabling the fast.html response to be delivered without delay, even while larger files (slow1/slow2) are in progress.

These nearly simultaneous timestamps indicate that all three streams were handled independently, and the small file fast.html was not delayed by the slower file downloads. This is because HTTP/3 runs over QUIC and UDP, which supports independent multiplexing per stream and avoids the TCP HOL blocking seen in HTTP/2.

9.

107	3.904461...	172.17.0.3	10.0.2.11	SCTP	84 INIT
108	3.904461...	172.17.0.3	10.0.2.11	SCTP	84 INIT
109	3.905356...	172.17.0.1	172.17.0.3	SCTP	356 INIT_ACK
110	3.905360...	172.17.0.1	172.17.0.3	SCTP	356 INIT_ACK
111	3.906427...	172.17.0.3	10.0.2.11	SCTP	312 COOKIE_ECHO
112	3.906427...	172.17.0.3	10.0.2.11	SCTP	312 COOKIE_ECHO
113	3.907067...	172.17.0.1	172.17.0.3	SCTP	52 COOKIE_ACK
114	3.907069...	172.17.0.1	172.17.0.3	SCTP	52 COOKIE_ACK
115	3.918593...	172.17.0.3	10.0.2.11	SCTP	104 DATA (TSN=0)
116	3.918593...	172.17.0.3	10.0.2.11	SCTP	104 DATA (TSN=0) (retransmission)
117	3.918636...	172.17.0.1	172.17.0.3	SCTP	64 SACK (Ack=0, Arwnd=106459)
118	3.918638...	172.17.0.1	172.17.0.3	SCTP	64 SACK (Ack=0, Arwnd=106459)
Frame 107: 84 bytes on wire (672 bits), 84 bytes captured on interface eth0, 84 bytes from 10.0.2.11 to 172.17.0.3 on interface eth0					0000 00 03 00 01 00 06 02 42 ac 11 00 03 00 00 08 00
Linux cooked capture v1					0010 45 02 00 44 00 00 40 00 40 84 82 15 ac 11 00 03
Internet Protocol Version 4, Src: 172.17.0.3, Dst: 10.0.2.11					0020 0a 00 02 0b eb 56 14 51 00 00 00 00 00 00 00 00
Stream Control Transmission Protocol, Src Port: 60246, Dst Port: 5201					0030 01 00 00 24 64 df 57 10 00 01 a0 00 00 0a ff ff
Source port: 60246					0040 ad a3 5e 08 00 0c 00 06 00 05 00 00 80 00 00 04
Destination port: 5201					0050 c0 00 00 04
Verification tag: 0x00000000					
[Association index: disabled (enable in preferences)]					
Checksum: 0x00000000 [unverified]					
[Checksum Status: Unverified]					
INIT chunk (Outbound streams: 10, inbound streams: 6)					

The SCTP handshake, as observed above, consists of four distinct steps:

INIT – The client begins the association by sending an INIT chunk.

INIT_ACK – The server responds with an INIT_ACK.

COOKIE_ECHO – The client then replies with a COOKIE_ECHO, echoing a cookie sent by the server in the INIT_ACK.

COOKIE_ACK – The server finalizes the handshake by sending a COOKIE_ACK.

This 4-way handshake is designed with a built-in cookie mechanism for protection against SYN flood attacks, unlike TCP. The client must return the server's cookie to prove its identity before the server allocates resources, making SCTP more resilient. By contrast, TCP (as used in HTTP/2) uses a 3-way handshake:

SYN – Client initiates the connection.

SYN-ACK – Server acknowledges and responds.

ACK – Client acknowledges back.

TCP's handshake is shorter but does not include a cookie verification mechanism, making it more vulnerable to spoofed connections and certain DoS attacks.

10. We don't see HEARTBEAT packets in the capture because they're only used when there's no data being sent for a while or when SCTP uses multiple paths (multi-homing). In our case, the connection was active with constant data transfer and only one path was used, so there was no need for HEARTBEAT messages.

11. QUIC, used in HTTP/3, runs over UDP, while SCTP is its own transport protocol built on IP. Both support sending multiple streams within a single connection, but SCTP

calls this "multistreaming" and QUIC calls it "multiplexing". QUIC includes built-in encryption and better NAT traversal, while SCTP is more focused on reliability and message boundaries. In terms of congestion control, both implement their own mechanisms, but QUIC is more optimized for modern web performance.